

Software Fault Isolation and Control Flow Integrity

Fritz Henglein

Lecture 6, Software Security, DIKU

2025-09-23

All slides by: Gang Tan, *Principles and Implementation Techniques of Software-Based Fault Isolation*, Penn State University, Spring 2019


Isolation via Protection Domains

2

- A fundamental idea in computer security
 - [Lamson 74] “Protection”
- Structure a computer system to have multiple **protection domains**
 - Each domain is given a set of privileges, according to its trustworthiness

Example: the Separation between OS and User Applications

3



- One OS domain (the kernel mode)
 - Privileged: execute privileged instrs; set up virtual memory; perform access control on resources; ...
- Multiple application domains
 - Go through OS syscalls to request access to privileged operations
 - Application domains are isolated by OS processes

Isolating Untrusted Components

4

- Using separate protection domains is a natural choice for isolating untrusted components
- E.g., isolating plug-ins in a web browser
 - Malfunctioning/malicious plug-ins must not crash or violate the security of the browser
- E.g., isolating device drivers in an OS

Many Forms of Protection Domains

5

- **Hardware-based virtualization:** Each domain in a virtual machine
 - Pros: easy to use; high degree of isolation
 - Cons: extremely high overhead when context switching between domains
- **OS processes:** each domain in a separate OS process
 - Pros: easy to use; cons: high context-switch overhead
- **Language-based isolation:** rely on safe languages or language features such as types
 - Pros: fine grained, portable, flexible, low overhead
 - Cons: high software engineering effort to use safe languages/features
 - Guaranteed safety and security by construction requires languages with effective support for reasoning about program semantics

Comparison of Forms of Protection Domains

6

	Context-switch overhead	Per-instruction overhead	Require compiler support	Software engineering effort
Virtual machines	Very high	None	No	None
OS processes	High	None	No	None
Language-based isolation	Low	Medium (dynamic checking) or none (static checking)	Yes	High
SFI	Low	Low	Maybe	None or medium

Per-instruction overhead: whether for each instruction additional checking is needed

Software-Based Fault Isolation (SFI)


7

- Introduced by [Wahbe et al. 93] for MIPS
 - PittSFleld [McCamant & Morrisett 06] extended it to x86
- SFI isolation is within the same process address space
 - Each protection domain has a designated memory region
 - Same process: avoiding costly context switches
- Implementation by inserting software checks before critical instructions
 - E.g., memory reads/writes, indirect branches.
- Pros: fine grained, flexible, low context-switch overhead
- Cons: may require some compiler support and software engineering effort

THE SFI POLICY

The SFI Sandbox Setup


9



- Data region (DR): [DB,DL]
 - Holds data: stack, heap
 - Data Begin, Data Limit
 - Data Limit may be an offset
- Code region (CR): [CB,CL]
 - Holds code
- Safe External (SE) addresses
 - Host trusted services that require higher privileges
 - Code can jump to them for accessing resources
- DR, CR, and SE are disjoint

The SFI Policy

10



Implications of the SFI Policy

11

□ Non-writable code

- All memory writes must write to DR
- Code region must not be modified
 - No self-modifying code

□ Non-executable data

- Control flow cannot transfer to the data region
- Cannot inject data to DR and execute it as code
 - Code injection prevented

Stronger Policies

12

- An SFI implementation might implement a stronger policy
 - For implementation convenience
 - For efficiency
- E.g., PittSFleld [McCamant & Morrisett 06]
 - Disallow jumping into the middle of instructions on x86, which has variable-sized instructions
- E.g., NaCl [Yee et al. 09]
 - Disallow system call instructions in the code region

SFI ENFORCEMENT OVERVIEW

SFI Enforcement Overview


14

- ***Dangerous instructions:*** memory reads, memory writes, control-transfer instructions
 - They have the potential of violating the SFI policy
- SFI enforcement
 - Checks every dangerous instruction to ensure it obeys the policy
- Two general enforcement strategies
 - Dynamic binary translation
 - Inlined reference monitors

Dynamic Binary Translation

15


- Efficient interpretation of instructions
- For a dangerous instruction, the interpreter checks it is safe according to the policy
- Examples
 - Program shepherding [Kiriansky et al. 02]
 - libdetox [Payer & Gross 11]
 - VX32 [Ford & Cox, 08]



Inlined Reference Monitors (IRM)

16

- A static program rewriter
 - Inlines checks into the input program
- More efficient
 - No dynamic translation costs
 - Can optimize checks via static analysis
- More trustworthy
 - A separate verifier can check that checks are inlined correctly
- The main SFI implementation strategy and the focus of the rest slides



Strategies for Implementing IRM Rewriters

17

- Binary Rewriting
 - Input: binary code
 - Steps: perform disassembly; insert checks; assemble the instrumented code
 - Pros: not requiring source code
 - Cons: hard to disassemble stripped binaries
- Inside a compiler
 - Input: source code
 - Steps: the compiler inlines checks when generating binary code
 - Pros: can perform more optimizations on checks with richer information on code (e.g., types)

ENFORCING SFI'S DATA-ACCESS POLICY AND OPTIMIZATIONS

An Idealized Assembly Language

19

- We introduce an idealized assembly language
 - For writing assembly-code examples to show SFI enforcement and optimizations

<i>(Instr)</i>	$i ::= r_d := r_s \ aop \ op$
	$ \quad r_d := \text{mem}(r_s + w) \quad \quad \text{mem}(r_d + w) := r_s$
	$ \quad \text{if } (r_s \ cop \ op) \text{ goto } w \quad \quad \text{jmp } op$
<i>(Register)</i>	$r ::= r0 \ \ r1 \ \ r2 \ \ ...$
<i>(Operand)</i>	$op ::= r \ \ w$
<i>(ALOp)</i>	$aop ::= + \ \ - \ \ \gg \ \ \ll \ \ \& \ \ '$
<i>(CompOp)</i>	$cop ::= > \ \ < \ \ \leq \ \ \geq \ \ = \ \ \neq \ \ ...$

* w for a static constant word

Abbreviations and Terminology

20

- $r := r' + 0$ abbreviated as $r := r'$
- In memory instructions, $\text{mem}(r+0)$ abbreviated as $\text{mem}(r)$
- Direct branches: $\text{jmp } w$
 - The jump target is a static constant word w
- Indirect branches: $\text{jmp } r$
 - The jump target is in a register and cannot always be statically determined

Example

21

```
r3 := r1
r4 := r2 * 4
r4 := r1 + r4
r5 := 0
loop:
if r3 ≥ r4 goto end
r6 := mem(r3)
r5 := r5 + r6
r3 := r3 + 4
jmp loop
end:
```


- r1 is a pointer to the beginning of an array
- r2 holds the array length
- The program computes in r5 the sum of array elements

Naïve Enforcement

22

- Insert checks before memory reads/writes

```
mem(r1+12) := r2 // unsafe mem write
```



```
r10 := r1 + 12  
if r10 < DB goto error  
if r10 > DL goto error  
mem(r10) := r2
```

*Assume r10 is a scratch register

Naïve Enforcement

23

- Sufficient for security
- Has a high runtime overhead
 - Two checks per memory access
- A practical SFI implementation
 - Need to implement a range of optimizations to drive down the cost
 - Discussed next
 - Side note: a good illustration of what's needed to make a simple security scheme practical

Optimization: Integrity-Only Isolation

24

- A program performs many more reads than writes
 - ▣ In SPEC2006, 50% of instructions perform some memory reads or writes; only 10% perform memory writes [Jaleel 2010]
- For integrity, check only memory writes
- Sufficient when confidentiality is not needed
- Much more efficient
 - ▣ [Wahbe et al. 1993] on MIPS using typical C benchmarks
 - 22% execution overhead when checking both reads and writes; 4% when checking only writes
 - ▣ PittSField on x32 using SPECint2K
 - 21% execution overhead when checking both reads and writes; 13% when checking only writes
- As a result, most SFI systems do not check reads

Optimization: Data Region Specialization

25

- Special bit patterns for addresses in DR
 - ▣ To make address checks more efficient
- One idea in the original SFI [Wahbe et al. 1993]
 - ▣ Data region addresses have the same upper bits, which are called the **data region ID**
 - ▣ Only one check is needed: check whether an address has the right region ID

Optimization: Data Region Specialization

26

- Example: DB = 0x12340000 ; DL = 0x1234FFFF
 - The data region ID is 0x1234
 - “mem(r1+12) := r2” becomes

```
r10 := r1 + 12
```

```
r11 := r10 >> 16 // right shift 16 bits to get the region ID
```

```
if r11 ≠ 0x1234 goto error
```

```
mem(r10) := r2
```

Q: What does “r3 := mem(r4-20)” become?

Optimization: Address Masking

27

- **Address checking** stops the program when the check fails
 - ▣ Strictly speaking, unnecessary for isolating faults
- A more efficient way: force the address of a memory operation to be a DR address and continue execution
 - ▣ Called **address masking**
 - ▣ “Ensure, don’t check”
 - ▣ When using data region specialization, just modify the upper bits in the address to be the region ID
 - ▣ PittSFeld reported 12% performance gain when using address masking instead of checking for SPECint2000

Optimization: Address Masking

28

- Example: DB = 0x12340000 ; DL = 0x1234FFFF
 - “mem(r1+12) := r2” becomes

```
r10 := r1 + 12  
r10 := r10 & 0x0000FFFF  
r10 := r10 | 0x12340000  
mem(r10) := r2
```

Force the address to
be in DR

Q: What does “r3 := mem(r4-20)” become?

Wait! What about Program Semantics?

29

- “Good” programs won’t get affected
 - “Good” programs won’t access memory outside DR
 - For bad programs, we don’t care about whether its semantics is destroyed
- Cons: does not pinpoint the policy-violating instruction
 - A downside for debugging and assigning blame

Optimization: One-Instruction Address Masking

30

- Idea
 - The data region ID has only a single bit on
 - E.g. 0x0001, 0x0002, 0x0004, 0x0008, 0x0010, ..., 0x1000, 0x2000, 0x4000, 0x8000 for 16-bit data region IDs
 - Make the zero-ID region, 0x0000 for 16-bit data regions, unmapped in the virtual address space
- A memory access is safe
 - If the address is either in the data region or in the zero-ID region
 - Reason: an access to the zero-ID region generates a hardware trap because it accesses unmapped memory
- Benefit: cut down one instruction for masking
 - PittSFleld reported 10% performance gain on SPECint2000

Optimization: One-Instruction Address Masking

31

- Example: DB = 0x20000000 ; DL = 0x2000FFFF
 - Region ID is 0x2000
 - “mem(r1+12):= r2” becomes

```
r10 := r1 + 12  
r10 := r10 & 0x2000FFFF  
mem(r10) := r2
```

- Result is an address in DR or in the (unmapped) zero-ID region
- Cons: limit the number of DRs
 - In a 32-bit system, if a DR's size is 2^n , then we can have at most $(32-n)$ DRs

Data Guards

32

- A **data guard** refers to either address checking or address masking
 - ▣ When which one is used is irrelevant
- Introduce a pseudo-instruction “ $r' = dGuard(r)$ ”
 - ▣ To hide implementation details
- An implementation should satisfy the following properties of “ $r' = dGuard(r)$ ”
 - ▣ If r is in DR, then r' should equal r
 - ▣ If r is outside DR, then
 - For address checking, an error state is reached
 - For address masking, r' gets an address within the safe range
 - The safe range is implementation specific; it's typically DR; for PittSField, it's DR plus the zero-ID region


Optimization: Guard Zones

33

- Place a guard zone directly before and after the DR
- First described by Wahbe et al. (1993); further extended by Zeng et al. (2001) and Sehr et al. (2010)

Guard Zones: Safe Accesses

34



- E.g., GSize=4k
- **Assumption:** Guard zones are unmapped
 - Thus, access to guard zones are trapped by hardware
- A memory read/write is **safe** if the address is in [DB-GSize, DL+GSize]

Guard Zones Enable More Optimizations

35

- In-place sandboxing
- Redundant check elimination
- Loop check hoisting

Similar to those optimizations performed in an optimizing compiler, enabled by classic static analysis

Optimization: In-Place Sandboxing

36

- A commonly used addressing mode in memory operations
 - A base register plus/minus a small constant offset
 - E.g., the register points to the start address of a struct, and the constant is the offset to a field
- In this case, just guard the base register in place is sufficient, when the constant is no greater than GSize

Optimization: In-Place Sandboxing

37

- Example: “ $\text{mem}(r1+12) := r2$ ” becomes

```
r1 := dGuard(r1)
```

```
mem(r1+12) := r2
```

- No need for a scratch register
- Why is the above safe?
 - “ $r1 := \text{dGuard}(r1)$ ” constrains $r1$ to be in DR and then $r1+12$ must be in $[\text{DB-GSize}, \text{DL+GSize}]$, assuming $\text{GSize} \geq 12$
 - Note: for PittSField, we need to have guard zones around the zero-ID region too, since dGuard constrains $r1$ to be either in DR or the zero-ID region in PittSField
 - Will ignore this for the rest of the slides

Optimization: In-Place Sandboxing

38

- NaCl-x86-64 (Sehr et al., 2010) implemented a similar optimization
- Put guard zones of 40GB above and below a 4GB sandbox
 - 64-bit machines have a large virtual address space
 - As a result, most addresses in memory operations can be guaranteed to stay in [DB-GSize, DL+GSize]
 - By carefully controlling the registers in “base register + a scaled index register + displacement” addressing mode

Optimization: Redundant Check Elimination

39

- Idea: perform range analysis to know the range of values of registers and use that to remove redundant data guards

```
r1 := dGuard(r1)           ← = = = = = = = = = = r1 ∈ [DB,DL]
r2 := mem(r1 + 4)
...
// r1 is not changed in between
r1 := dGuard(r1)           ← = = = = = = = = = = r1 ∈ [DB,DL]
r3 := mem(r1 + 8)
```

Removing the redundant guard

Optimization: Loop Check Hoisting

40

- Idea: a guard in a loop is hoisted outside
 - The guard is performed only once per loop instead of once per loop iteration
- Key observation
 - If $\text{addr} \in [\text{DB-GSize}, \text{DL+GSize}]$, then a successful (untrapped) memory operation via addr means $\text{addr} \in [\text{DB}, \text{DL}]$
 - If it were in any of the guard zones, then a trap would be generated

Loop Check Hoisting Example

41

Before optimization

```
r3 := r1  
r4 := r2 * 4  
r4 := r1 + r4  
r5 := 0  
loop:  
    if r3 ≥ r4 goto end  
    r3 := dGuard(r3)  
    r6 := mem(r3)  
    r5 := r5 + r6  
    r3 := r3 + 4  
    jmp loop  
end:
```

After optimization

```
r3 := r1  
r4 := r2 * 4  
r4 := r1 + r4  
r5 := 0  
r3 := dGuard(r3)  
loop:  
    if r3 ≥ r4 goto end  
    r6 := mem(r3)  
    r5 := r5 + r6  
    r3 := r3 + 4  
    jmp loop  
end:
```

* r1 is a pointer to the beginning of an array; r2 holds the array length; the program computes in r5 the sum of array elements

Why is the Previous Optimized Code Safe?

42

```
r3 := r1
r4 := r2 * 4
r4 := r1 + r4
r5 := 0
r3 := dGuard(r3)
loop:
    if r3 ≥ r4 goto end
    r6 := mem(r3)
    r5 := r5 + r6
    r3 := r3 + 4
    jmp loop
end:
```

← = = = = r3 ∈ [DB,DL]
← = = = = r3 ∈ [DB,DL+4]
← = = = = r3 ∈ [DB,DL+4]
← = = = = r3 ∈ [DB,DL+4]
← = = = = r3 ∈ [DB+4,DL+4]

[DB, DL+4]
 \subseteq [DB-GSize, DL+GSize]

Optimization: Guard Changes Instead of Uses

43

- Some registers are used often
 - E.g., in 32-bit code, ebp is usually set in the function prologue and used often in the function body
- Idea
 - Sandbox the changes to those special registers, instead of uses
 - E.g., `ebp := esp` becomes
 - `ebp := esp`
 - `ebp := dGuard(ebp)`later uses of `%ebp` plus a small constant do not need to be guarded, if used together with guard zones

Scratch Registers

44

- The SFI rewriting may require finding scratch registers to store intermediate results
 - ▣ E.g., r10 in many of our previous examples
- If the old values of scratch registers need to be preserved
 - ▣ Need to save and restore the old values on the stack
- How to avoid that?

Optimization: Finding Scratch Registers

45

- Binary rewriting
 - Perform binary-level liveness analysis to find dead registers as scratch registers [Zeng et al. 11]
- Compile-level rewriting
 - Approach 1: reserve dedicated registers as scratch registers
 - E.g., PittSField reserves ebx as the scratch register by passing GCC a special option
 - Downside: increase register pressure
 - Approach 2: rewrite at the level of an IR that has unlimited number of variables
 - E.g., LLVM IR
 - A later register allocation phase maps those variables to registers or stack slots

Architecture-Specific Optimization

46

- An SFI implementation can use specific hardware features for efficient sandboxing
- NaCl and VX32 on Intel x32
 - Use x32's segmentation support
 - Data segment: base gets DB and limit and DL
 - Hardware automatically performs checks
 - However, not supported in x64
- ISBoxing On x64 [Deng et al. 15]
 - Put the data region in the first 4GB
 - Add address-override prefix to a memory instruction
 - Cons: only support one data region with a fixed size
- ARMlock on ARM [Zhou et al. 14]
 - Use ARM's memory domain feature

ENFORCING SFI'S CONTROL-FLOW POLICY

Control-Flow Policy

48

- Recall the policy: control-flow targets must stay in $[CB, CL] \cup SE$
- However, when using the IRM approach for SFI enforcement
 - Must also restrict the control flow to disallow bypassing of guards

Risk of Indirect Branches

49

I1: $r10 := r1 + 12$

I2: $r10 := \text{dGuard}(r10)$

I3: $\text{mem}(r10) := r2$

- Worry: what if there is a return instruction somewhere else and the attacker corrupts the return address so that the return jumps to I3 directly?
 - Then the attacker bypasses the guard at I2!
 - If attacker can further control the value in r10, then he can write to arbitrary memory location

Risk of Indirect Branches

50

- In general, any **indirect branch** might cause such a worry
 - If not carefully checked, it may bypass the guard
- Indirect branches include
 - Indirect calls (calls via register or memory operands)
 - Indirect jumps (jumps via register or memory operands)
 - Return instructions
- In contrast, direct branches are easy to deal with
 - Targets of a direct branch encoded in the instruction; can statically inspect the target

The Original SFI Solution [Wahbe et al. 93]

51

- Make r10 (in MIPS) a dedicated register
 - r10 only used in the monitor code, not used by application code
 - Also maintain the invariant that r10 always contains an address in DR before any branch
 - So even if an indirect branch bypasses the guard before a memory operation, the memory access stays within DR
- Cons?
 - Reduce the number of registers available to application code
 - Allow an indirect branch to target the middle of an instruction; problem for variable-sized instruction sets

A More Direct Approach: Control-Flow Integrity

52

- Define a **pseudo-instruction**
 - Either a non-dangerous instruction
 - Or a guard followed by a dangerous instruction
- **Strengthened control-flow policy**
 - All control-flow transfers must target the beginning of a pseudo-instruction in CR or an address in SE
- Note the strengthened policy rules out
 - Bypassing a guard
 - And jumping into the middle of an instruction

Aligned-Chunk Enforcement (PittSFleld)

53

- Divide the code into chunks of some size
 - E.g., 16 or 32 bytes
- Each chunk starts at an aligned address
 - addr is aligned if $\text{addr} \bmod \text{chunkSz} = 0$
- Make dangerous instrs and their guards stay within one chunk
 - E.g., “`r10 := dGuard(r10); mem(r10) := r2`” stay within one chunk
- Insert guards before indirect branches so that they target only aligned addresses (chunk beginnings)

Example

54

- Assume
 - CR is [0x10000000, 0x1000FFFF], the one-bit code region 0x1000
 - Chunk size is 16 bytes
 - Zero-ID region [0x00000000, 0x0000FFFF] unmapped
- Then “jmp r” becomes

$r := r \& 0x1000FFF0$


`jmp r`

- Q: why does the above ensures that the target address is (1) in CR or zero-ID region, and (2) a chunk beginning
- after &, r's upper 16 bits must be either 0x0000 or 0x1000
 - after &, r's lower four bits must all be 0, meaning it's 16-byte aligned

Downside of Aligned-Chunk Enforcement

55

- All legitimate jump targets have to be aligned
 - No-ops have to be inserted for that



- Extra no-ops slow down execution and increase code size
 - In PittSField, inserted no-ops account for half of the runtime overhead; NaCl-JIT incurs 37% slowdown because of no-ops
 - In NaCl-x64, the code size becomes 60% larger

Bitmap Based Enforcement (MIP [Niu & Tan, 13])

56

- Allow variable-sized chunks
 - A guard and the following dangerous instr still stay within one chunk
 - Chunk beginnings are remembered in an immutable bitmap
 - $b[\text{addr}] = 1$ iff addr is the beginning of a chunk
 - Before an indirect branch, insert a guard to check if $b[\text{addr}]$ is 1, assuming addr is the target
 - If not, jump to error
- Benefit: no need to insert no-ops
 - MIP-x32: 4% runtime overhead; 13% code increase
 - MIP-x64: 7% runtime overhead; 16% code increase

Fine-Grained CFI

57

- Enforce that a program follows a fine-grained control-flow graph
 - ▣ [Abadi et al, 05] and many other follow-up work
 - ▣ E.g., for each return, the fine-grained CFG defines a set of possible return targets
- Stronger than the pseudo-instruction based CFI policy
- Pros: we can use the fine-grained CFI to optimize away more guards [Zeng et al., 11]
- Cons: enforcing it incurs additional overhead
 - ▣ Unnecessary for the control-flow policy in SFI


Jumping Outside of Fault Domains

58

- Total isolation is rarely what's desired in practice
- Sandboxed code must interact with other parts of the system for its functionality
 - ▣ E.g., a browser plug-in must communicate with the browser's core for exchanging data with the core and other plug-ins

Allow Only Controlled Interaction

59



- The sandboxed code can jump to a pre-defined set of SE (Safe External) addresses
- Each SE address holds a trusted service
 - E.g., service for invoking OS syscalls (fopen, fread, ...)
 - E.g., service for allowing communication with other fault domains

Trusted Services

60

- Implemented outside of the fault domain
- They can implement additional security policies
 - E.g., can restrict fopen to open files only in a particular directory
 - Or can disallow fopen completely
 - Just do not set up a service entry for fopen

APPLICATIONS OF SFI

SFI Applications Overview

62

- Isolating OS kernel modules such as device drivers
 - MiSFIT [Small 97]; XFI [Erlingsson et al. 06]; BGI [Castro et al. 09]; LXFI [Mao et al. 11]
- **Isolating plug-ins in Chrome**
 - NaCl [Yee et al. 09]; NaCl-x64 [Sehr et al. 10]
- **Isolating native libraries in the Java Virtual Machine**
 - Robusta [Siefers et al. 10]; Arabica [Sun & Tan 12]

Google's Native Client (NaCl)

63

- SFI service in Chrome
 - [Yee et al. Oakland 09]
- Goal: download native code and run it safely in the Chrome browser
 - Much safer than ActiveX controls
 - Much better performance than JavaScript, Java, etc.
- Google's main motivation: run native-code games in Chrome



DOOM in NaCl

NaCl: Code Verification

64

- Code is verified before running
 - Allow restricted subset of x86 instructions
 - No unsafe instructions: memory-dependent jmp and call, privileged instructions, modifications of segment state, ...
 - Ensure SFI checks are correctly implemented for the SFI policy

NaCl Sandboxing

65

- x86-32 sandboxing based on hardware segments
 - Sandboxing reads and writes for free
 - 5% overhead for SPEC2000 benchmarks
- However, hardware segments not available in x86-64 or ARM
 - Use instructions for address masking [Sehr et al. 10]
 - x86-64/ARM: 20% for sandboxing mem writes and computed jumps

NaCl SDK

66

- Modified GCC tool-chain
 - Inserts appropriate masks, alignment requirements
- Trampolines allow restricted system-call interface and also interaction with the browser
 - Pepper API: access to the browser, DOM, 3D acceleration, etc.


Robusta [Siefers, Tan, Morrisett 10]

67

- SFI service in a Java Virtual Machine (JVM)
 - Allow Java code to invoke native code safely through the Java Native Interface (JNI)
- The basic idea
 - Put native code in an SFI sandbox and allows only controlled access to JVM services

Robusta [Siefers, Tan, Morrisett 10]

68



Native Code Threat

- ❑ Direct JVM mem access
- ❑ Abusive JNI calls
- ❑ OS syscalls

Robusta Remedy

- ❑ SFI: Prevent direct JVM access
- ❑ Perform JNI safety checking
- ❑ Reroute syscall requests to Java's security manager

FUTURE DIRECTIONS

Future Directions

70

- Tool and programming support for program partitioning
 - ▣ How to turn a monolithic application into components in separate protection domains?
 - Privilege separation
 - ▣ It took Google significant effort to privilege separate Chrome into a system of cooperating processes [Barth et al. 08]

Future Directions

71

- Security enforcement on interface code
 - Trusted services in SE addresses are security critical
 - Experience shows that bugs are plenty in such interface code
 - Should apply program analysis/verification for bug finding
 - Or take a specification about interface security and enforce the security a la LXFI [Mao et al. 11]

Future Directions

72

- Side channel control
 - SFI provides memory isolation but side channels are possible
 - E.g., we might structure a server to have a trusted core and have a sandbox to handle each client connections
 - However, if the core maintains some state that is shared by all connections, there might be a side channel
 - Similar channels were discovered in TCP (“Off-Path TCP Exploits: Global Rate Limit Considered Dangerous”)

Future Directions

73

- Recovery mechanism
 - Address checking terminates the sandbox when there is an illegal access
 - May still need to release resources
 - Address masking turns an illegal access to a legal one
 - May cause a benign but buggy sandboxed component to misbehave
 - It does not pinpoint the violating instruction
 - [Seltzer et al. 96] Wrap sandbox calls in transactions
 - Transactions are aborted when sandbox misbehavior is detected; resources are released as a result

More in the Survey Article

74

- G. Tan “Principles and Implementation Techniques of Software-Based Fault Isolation”, Foundations and Trends in Privacy and Security: Vol. 1, No. 3, pp 137–198.
 - <http://www.cse.psu.edu/~gxt29/papers/sfi-final.pdf>
- SFI verifier
 - Verifies that the result after SFI rewriting is correct
 - Basic idea and formalization
- References